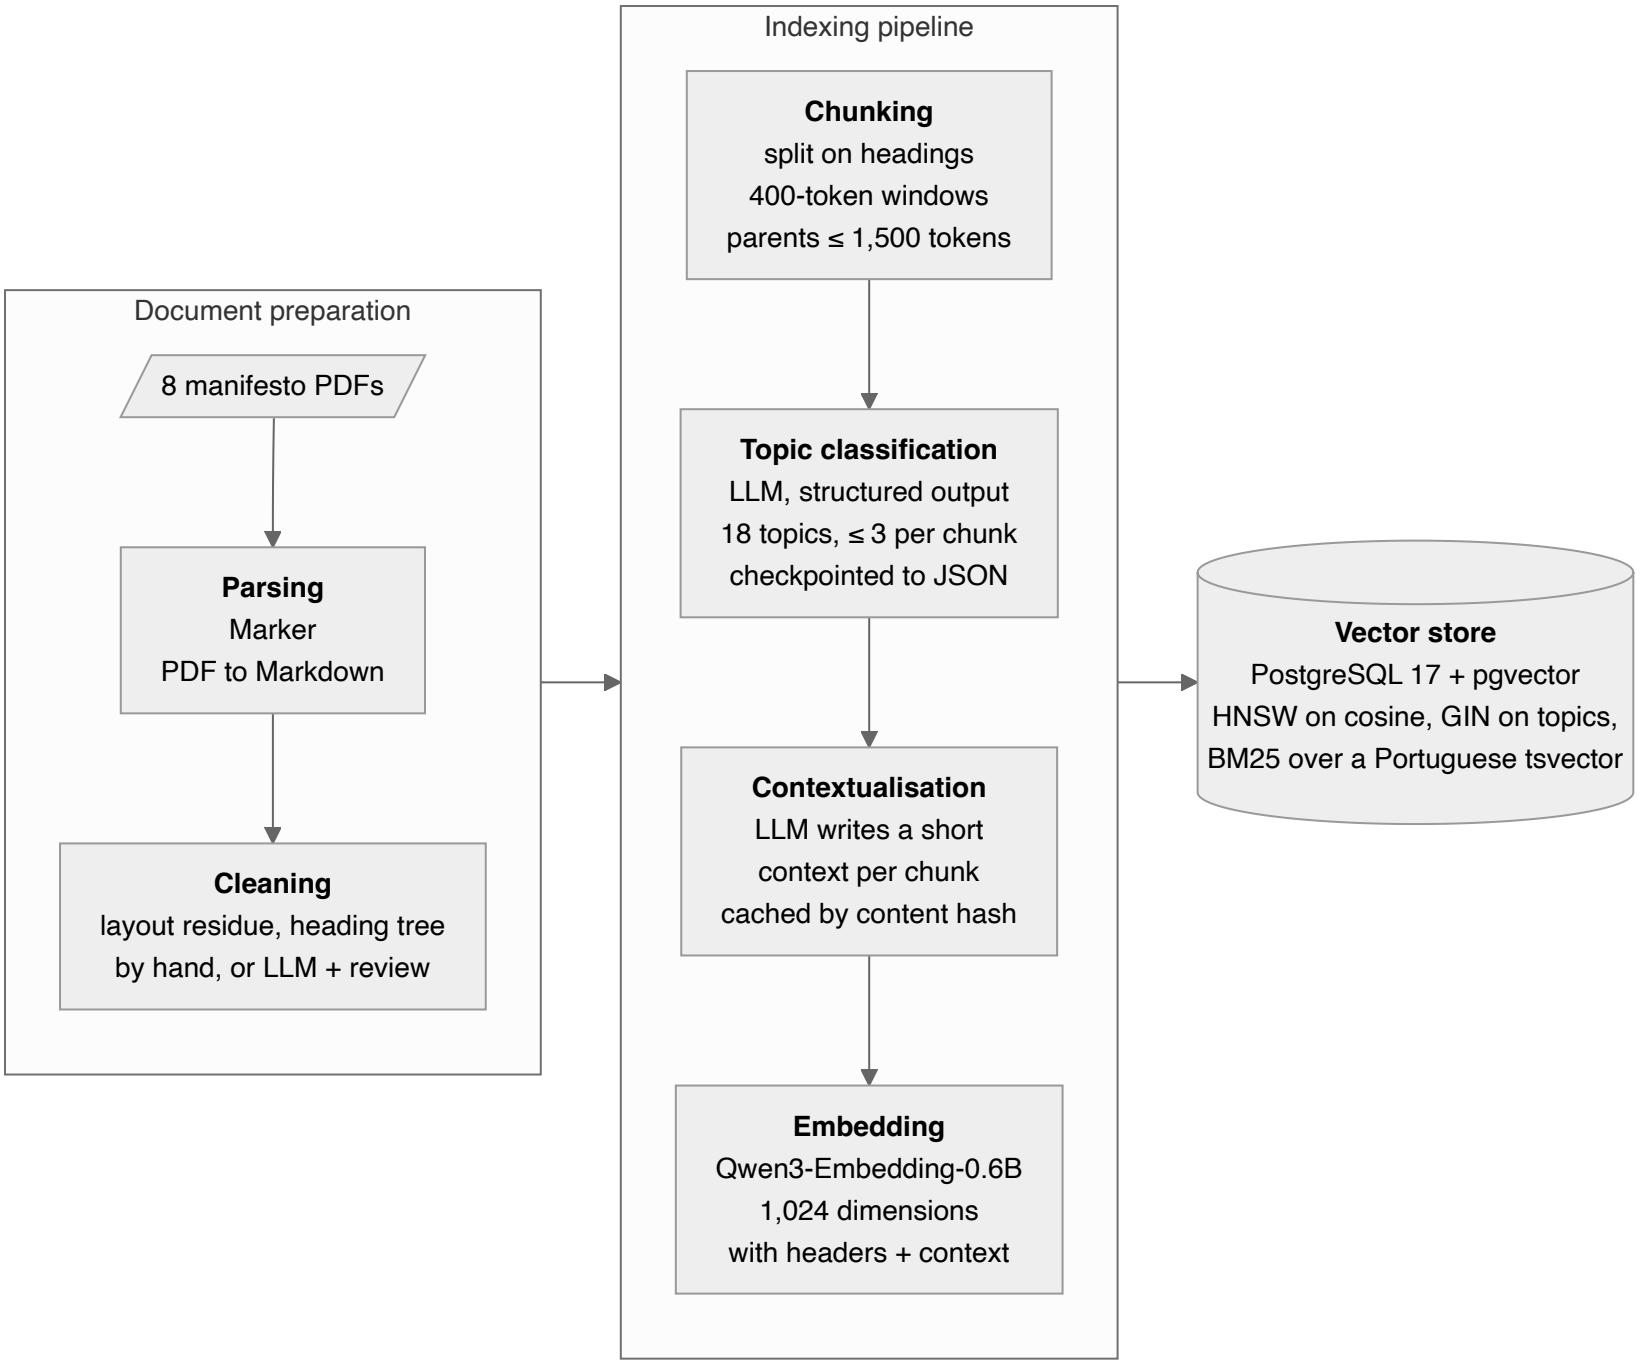
Document preparation
8 manifesto PDFs
Parsing
Marker
PDF to Markdown
Cleaning
layout residue, heading tree
by hand, or LLM + review
Indexing pipeline
Chunking
split on headings
400-token windows
parents ≤ 1,500 tokens
Topic classification
LLM, structured output
18 topics, ≤ 3 per chunk
checkpointed to JSON
Contextualisation
LLM writes a short
context per chunk
cached by content hash
Embedding
Qwen3-Embedding-0.6B
1,024 dimensions
with headers + context
Vector store
PostgreSQL 17 + pgvector
HNSW on cosine, GIN on topics,
BM25 over a Portuguese tsvector